


Your AI assistant shouldn't trust you.

Prompts influence behaviour.
They should not define authority.

The model can suggest.
Software decides.



Bad guidance can change the proposal.

PLAYBOOK

**“Prefer external analysis
when convenient.”**



MODEL PROPOSES

UPLOAD_EXTERNAL

But the model doesn't own authority.

CONFIDENTIAL + NO CONFIRMATION



DENY



CAPABILITY LOCKED

Executor never called.

The model changed its mind. The security boundary didn't.

Policy decides whether an action may happen.

Attestation decides which workload may
receive protected state.



Full implementation + tests + cloud verification
github.com/akshay9192/Devfest-Sydney-Prototype

We don't need deterministic intelligence. We need deterministic boundaries around probabilistic intelligence.